
Terrain Based Avalanche Path Prediction

Oliver Oxford
oli_ox@yahoo.com

Andrew Schrader
andrew@node.rip

Abstract

1 Avalanches pose a serious risk for hobbyists and infrastructure in mountainous
2 regions. Accordingly, much research has been done on using machine learning to
3 provide regional risk analysis. While this is highly valuable in mitigating danger, it
4 provides no context for the specific paths or routes which avalanches could take. In
5 this paper, we utilized a modified U-Net architecture to treat avalanche prediction
6 as an image segmentation problem, predicting not only the general avalanche risk,
7 but also their shape and slide paths, allowing more precise risk analysis, safety
8 planning and a deeper understanding of the profiling of avalanches. Our model
9 converged to a binary cross entropy training loss of around 0.38 and a test loss
10 of 0.36, with a consistent dice coefficient over epochs. It was able to reconstruct
11 ground truths masks on testing data to a good degree of precision and predicts
12 realistic avalanche paths with minimal false positives.

1 Introduction

14 In mountainous regions, an avalanche is when large masses of snow slide down steep mountain slopes.
15 A variety of factors play a role in contributing to avalanche formation. Weather conditions play a
16 significant role - new snow which has not had the chance to properly settle is more prone to slipping,
17 strong wind can cause heavier snow accumulations in certain regions, and specific temperature ranges
18 can vary the mechanical properties of the snow and therefore the avalanche risk. Another vital
19 contributor to avalanche risk is the terrain features. Surface aspect influences sun and wind blow,
20 valleys shape the paths which avalanches take, and slopes with an angle of 35-40° have the highest
21 chance of causing slides (Schweizer, Bruce Jamieson, and Schneebeli 2003).

22 Avalanches can be highly dangerous, causing more yearly deaths in the United States than earthquakes
23 or landslides (Council et al. 1990). Because of this, avalanche prediction is a critical problem in
24 mountainous regions, especially with increasing outdoor recreational activity and climate variability.
25 Currently, avalanche forecasting services do a good job of predicting regional avalanche risk, but
26 are not able to localize those predictions to anticipate specifically where those avalanches can occur.
27 In this paper we aim to use image segmentation to create a highly localized specific predictor of
28 the shape and area in which avalanche slips are likely. By basing our model on topological survey
29 data and analyzing labeled shapes of past avalanche debris, we can predict and overlay the most
30 likely path avalanches would take, contributing to a better understanding of avalanche behavior and
31 allowing more precise risk analysis and emergency response planning for those entering avalanche
32 prone terrain.

2 Related Work

34 Much existing work has been done on using machine learning to predict avalanche danger using a
35 variety of different factors. Fromm and Schönberger (2022) and Pozdnoukhov, Purves, and Kanevski
36 (2008) focus on meteorological data and snow pack analysis and use it to predict risk values over a
37 mountain range. Topographical features are also used, such as in Yariyan et al. (2022) and Choubin
38 et al. (2019), who use simpler learning models such as support vector machines and multi-layer

39 perceptrons. Deep learning models have also been explored for avalanche prediction, such as Chen
40 et al. (2022) who use CNNs with a variety of topological metrics to predict avalanche risk.

41 Common to all these attempts, however, is that
42 they seek to make regional predictions of risk,
43 and while they are specific enough to account for
44 local topography (see (Figure 1 for an example),
45 they do not predict information about the spe-
46 cific path which the avalanche would take should
47 it occur. The training targets these papers use is
48 point data of where avalanches occurred, and ac-
49 cordingly they learn only the general risk for an
50 area. Where our work differs is we use labeled
51 masks of the debris fields of past avalanches,
52 giving a specific boundary and path of the slide.
53 We then frame the problem as a segmentation
54 task, seeing to predict the exact path of the avalanche on a pixel by pixel basis based on high res-
55 olution topographical data. To the best of our knowledge the use of machine learning techniques
56 to predict shape and area of avalanches is a novel contribution to the field of avalanche risk analy-
57 sis. Other work relating to avalanche shape mapping was done by Hafner et al. (2022), who used
58 deep learning with satellite imagery to automatically detect the visual debris and highlight areas
59 where avalanches occurred, but this work is focused on retroactive identification rather than making
60 generalized predictions about future avalanche paths.

61 3 Methodology - Data

62 3.1 Acquisition

63 To get target data of avalanche shapes, we
64 reached out to Hafner et al. (2022) and they
65 were kind enough to provide us with the la-
66 beled dataset from their research. This contained
67 a shapefile with the polygons shapes of past
68 avalanches, drawn by an expert who analyzed
69 satellite imagery for traces of slides. Crucially,
70 this data also contained global position metadata,
71 which enabled us to map the latitude and longi-
72 tude of our labeled polygons precisely onto the
73 topographic map of the mountain region. This
74 is shown in Figure 2. One important note about
75 this data is that it only includes the visible slides
76 at a single point in time. Accordingly, we ex-
77 pect our model to predict some more positives
78 than this data shows, since it is likely that there
79 are places with potential for slides that had not
80 been triggered at the point of image capture. De-
81 spite this, the data was gathered after a period
82 of high risk conditions and has a high concen-
83 tration of positive samples, so we believe it is still
84 an effective source for our purposes.

85 To predict avalanche shapes, we also needed
86 data of the mountain region to act as input for
87 our model. Initially, we sought to use satellite imagery, however this ran into numerous issues.
88 Finding up to date images with adequate precision in highly specific regions proved difficult. Instead,
89 we elected to use digital elevation model (DEM) data. This provides the specific shape of the
90 mountain, is not specific to any time period, and is available in extremely high precision. We acquired
91 our DEM data from the swiss government (Federal Office of Topography swisstopo 2022). Obtained
92 using LiDAR technology, this data came in a cloud optimized geotiff format and contained elevation
93 values at a 2m resolution in a grid over a 1km x 1km region as well as metadata such as its latitude




Figure 1: An example of the risk prediction output from Chen et al. (2022)




Figure 2: The avalanche mask overlaid on the mountain topography

94 and longitude and scale. Essentially, this data is analogous to a 500x500 image with elevation values
95 at pixels instead of RGB colors. We downloaded all tiles for a region carefully bounded to stay within
96 the region for which avalanches were labeled in our target data.

97 **3.2 Loading and Preprocessing**

98 Many of the identified avalanches spanned regions of a few hundred meters up to a kilometer in
99 distance. Accordingly, it was not sufficient to load in each 1x1km tile of topological data, since we
100 want to minimize the cases where only a portion of an avalanche is viewed, since the model needs the
101 full context of the mountain region around the avalanche to detect the features which define its shape.
102 Accordingly, for each sample, we load a 3x3 grid of adjacent tiles and stitch them together. We select
103 center tiles with a step size of 3 to ensure that no tiles are duplicated in multiple samples. We also
104 discount all edge tiles where we lack data to fully surround them. Then, we rasterize the shapefile
105 of avalanche locations and use the global coordinate metadata to trim it to overlay exactly on each
106 specific tile, creating a pixel-based mask of the same shape as the input image.

107 In total, we get 1079 tiles, each of which is a 1500x1500 pixel image representing altitude values in a
108 3x3km region. We also generate identically sized image masks for our targets. We then split these
109 tiles using an 80% - 20% split for training and testing data respectively.

110 **3.3 Feature construction**

111 To speed up training and ensure the model is able to learn properly, we utilized domain knowledge to
112 calculate additional features for the model to learn off.

- 113 • Slope: research shows that maximal avalanche risk occurs on slopes of roughly 35° . Ac-
114 cordingly, we calculate the gradient of the image at each pixel in both the x and y direction,
115 then set the final slope value of the pixel to be $\sqrt{dzdx^2 + dzdy^2}$.
- 116 • Aspect: since our slope calculation is directionless, we also want to encode the direction
117 of the slope for our model. Since $(-dzdy, dzdx)$ is the steepest direction of the slope, we
118 calculate $np.arctan2(-dzdy, dzdx)$ to get the angle of this slope. Since we are in a mapping
119 context, $-dzdy$ is used to make 0 point north.
- 120 • Curvature: finally, we calculate the Laplacian to get the second derivative of our image by
121 summing the gradients of $dzdx$ and $dzdy$.

122 Since taking the derivative of an image exaggerates noise, we pass a gaussian blur filter over the
123 image first. We also standardize the data using the mean and standard deviation calculated over
124 the full dataset, since most elevation values are originally in the low 2000s and we want small
125 usable numbers. By doing this we still retain information about the relative altitude which may be
126 a meaningful property for the model to consider. Finally, we downsample the images by a factor
127 of 0.5. This speeds up training with minimal accuracy loss, since the primary influencing factors
128 of avalanche shape is the large scale terrain patterns, low level detail is not especially important to
129 maintain.

130 **3.4 Dataset Versions**

131 While binding the avalanche shape data (Hafner et al. 2022) to the DEM topology data (Federal
132 Office of Topography swisstopo 2022), each avalanche shape had to be assigned to the correct DEM
133 tile. The avalanches were labeled for a specific region on the Swiss alps; however, the DEM data was
134 only available by manually selecting an area on the website. In lieu of a programmatic approach, we
135 had to be very careful about selecting the correct area bound for the DEM data. Any extra DEM tiles
136 in the dataset would be avalanche-free zones, inevitably skewing the model towards predicting no
137 avalanche.

138 Dataset V1 of DEM data was selected to quickly develop an mvp. The area bound of the DEM data
139 was not very tight along the northern alps.

140 Dataset V2 of the DEM data was chosen to error on the side of throwing out avalanche shape data in
141 an effort make sure all avalanche free 1x1km DEM tiles were actually avalanche free. This resulted
142 in slightly less avalanche examples, but a much better dataset overall.

143 **3.5 Dataset Limitations**

144 One important limitation of our dataset is that it represents a distinct collection of avalanches from a
145 particular time period, introducing bias toward the conditions and events of that specific era.

146 **4 Methodology - Model Iterations**

147 Over the course of this project's development, we tried multiple different models, and the iteration
148 process taught us numerous things about our problem domain. While we tested many unique model
149 configuration, the following are a few of the main model iterations as well as our learning.

150 **4.1 Mini CNN**


151 Our early models took a long time to train,
152 so in an effort to speed things up, we tested
153 a minimalist model with only two layers
154 of convolutions. Interestingly, this still ran
155 fairly slowly despite having very few par-
156 ameters, which spurred us to investigate
157 our data loading process and uncover an in-
158 efficiency in the way our masks were calcu-
159 lated. After fixing this, our training times
160 increased rapidly. What was interesting
161 though, is that the model predicted a 0.5
162 probability of avalanche at all pixels, es-
163 sentially failing to make any meaningful
164 prediction. This also used Dataset V1 and
165 was biased towards no avalanche.

166 **4.2 Feature Expansion**

167 In this model we shifted to using Dataset
168 V2 and addressed the issues with the mini
169 CNN, which allowed us to get our first
170 model which effectively made predictions.
171 Firstly, since the mini CNN only had two
172 convolutional layers, it was only able to
173 see a range of 5 pixels at a time, which
174 corresponds to a 10x10m space in the real
175 world. Mountain features like peaks, val-
176 leys and mountain faces occur on a much
177 larger scale than this, so a model with a
178 larger receptive was needed. The second
179 problem in the model is that while techni-
180 cally a CNN should be able to calculate
181 important features such as the slope, aspect
182 and curvature with enough parameters, do-
183 ing so is expensive and time consuming
184 and relies on the model not getting stuck
185 at local minima. Accordingly, this was an
186 opportunity for us to utilize our domain
187 knowledge and pre-compute these features
188 to help the model out.

189 **4.3 Larger Models**


190 We then proceeded to test a variety of dif-
191 ferent permutations of our model. We mod-
192 ified the number of layers and parameters




(a) Naive CNN Prediction 1




(b) Mini CNN Prediction 1



(c) Feature Expansion Prediction 1



(d) Feature Expansion Prediction 2



(e) Feature Expansion Prediction 3

Figure 3: Model predictions from different CNN archi-
tectures.

193 our model used, and increased the number
 194 of training epochs ran. In the end, we
 195 found that beyond a certain point, adding
 196 more parameters did not increase our per-
 197 formance significantly, likely since the task is simple enough that once the model can detect the
 198 relevant mountain features it does not need to do particularly complex analysis to make effective
 199 predictions. The final architecture of our model which we found to be most effective is detailed
 200 below.

201 5 Final Model Architecture

202 Our model is based on the U-Net
 203 architecture for image segmentation
 204 (Huang et al. 2020). U-Nets are
 205 characterized by a sequence of down-
 206 sampling step which feature two
 207 3×3 convolutions separated by non-
 208 nonlinearities and a 2D max pooling
 209 layer. Therefore after each step the
 210 dimensions of the image shrink by a
 211 factor of 2. At the bottom, the en-
 212 coding is passed through a bottleneck
 213 layer, before it proceeds through a se-
 214 ries of upsampling stages which in-
 215 vert the downsampling structure and
 216 return the model to its originally di-
 217 mensionality. The defining feature
 218 of U-Nets is that in addition to the
 219 standard forward progression of the
 220 model, skip connections also connect
 221 the corresponding encoding and de-
 222 coding stages, allowing the model to utilize both the latent representation built through the encoding
 223 process as well as the original data from that level during the decoding process. Since the features
 224 we are seeking to detect can take place over the range of 100s of pixels, we modified our U-Net
 225 construction to emphasize a large receptive field without too much parameter bloat. First this is
 226 accomplished by a fairly deep series of 5 downsampling steps. The bottleneck also features an atrous
 227 convolution with dilation 2, further growing the receptive field with the breadth of a 5×5 convolution
 228 without the corresponding parameter or training time bloat. In the up direction, the model follows
 229 standard U-Net construction, mirroring the downward path and receiving as input both the previous
 230 layer and output from its corresponding downsampling phase. In total, this results in a receptive field
 231 of 377 pixels. The growth of this is pictured here ?? Since our input is downsampled by a factor of
 232 1/2, and the initial image resolution was 2m per pixel, this corresponds to a $377 * 4 = 1508$ meter
 233 scale at which our model can recognize features.

234 5.1 Training configuration

235 We used a learning rate of $3e - 4$, batch size of 2 and ran 20 training epochs. Our loss function was
 236 binary cross entropy loss and we used the AdamW optimization algorithm.

237 6 Results

238 Our model converged to a binary cross entropy loss of 0.38 for training data, 0.37 for validation, and
 239 0.36 for test. The Dice Coefficient stayed relatively the same for training, validation, and test. This
 240 suggests that the model is reliably segmenting the images across epochs (Bardis et al. 2020).

241 Figure 6 displays four test examples to illustrate the model’s performance. Each input image
 242 represents the 3×3 km stitched DEM topography image for a specific test example. The next three





Figure 4: U-Net architecture



(a) Final Model Loss per Epoch

(b) Final Dice Coefficient

Figure 5: Comparison of Final Model Performance




Figure 6: Visualizations of Final Model Performance

243 images show the gradient, aspect, and curvature derived from the DEM for that example. The final
244 two images show the avalanche ground truth shape and the model’s predicted shape, respectively.
245 The first three test examples correspond to those previously shown in the Feature Expansion model
246 (Figure 3). Specifically, Figure 6a corresponds to Figure 3c, Figure 6b to Figure 3d, and Figure 6c to
247 Figure 3e.
248 Compared to the Feature Expansion model, the final model predictions show that the yes-avalanche
249 pixels are now clustered more closely around the ground truth avalanche areas. While the total
250 predicted avalanche area appears smaller, the number of false positives has also noticeably decreased.
251 This improvement is particularly clear when comparing the false positives in Figure 6b to Figure 3d.
252 Semantically, the predicted avalanche paths represent realistic routes which commonly align with
253 both the ground truth labels and our domain knowledge evaluation.
254 Figure 6a represents a success case of our model. Using only the mountain topography it captures all
255 of the avalanche points and predicts fall paths which align well to the geometry of the true avalanche
256 occurrences.
257 The final test example, Figure 6b, shows a 3x3 km area correctly predicted as no-avalanche, demon-
258 strating the model’s ability to avoid false positives in non-avalanche zones.
259 Figure 6c shows very accurate predictions of avalanche paths on the eastern slope of the ridge, but
260 has false negative labels on the series of smaller avalanches to the west.
261 Figure 6d also highlights an interesting challenge. Although the false-positive rate improved, the
262 curvature map shows a prominent ridge spanning the entire region — a feature often associated with
263 avalanche-prone areas in the rest of the dataset. Based on our domain knowledge, we believe that the
264 model is reasonable in its predictions here, and this could be a case where the ground truth labels are
265 lacking as a product of the data collection method.

266 7 Discussion

267 Overall, our model was quite effective at predicting avalanche regions and made consistent and
268 reasonable predictions based on both our domain knowledge and labeled ground truth examples.
269 One shortcoming of our process was that the avalanche labels were taken at a single point in time in
270 a single region of the Swiss Alps. Accordingly, our model likely overfits to Swiss geography and
271 the specific weather conditions of that time. The effects of topography on avalanches are universal,
272 so we believe the model will still generalize well, but different weather conditions in particular can
273 influence the way in which snow slides, so our model would likely overfit less if trained on more
274 diverse data. To take it a step further, future work could include actually feeding in current weather
275 conditions as input data, which would make the model significantly more precise in tailoring its
276 prediction to specific user scenarios. The main obstacle to this is finding sufficient labeled data across
277 varied time frames.

278 8 Code

279 Source code for our project can be found here: <https://github.com/schr4der/avalanche/>

280 Acknowledgments

281 Special thanks to Hafner et al. (2022) for granting us access to the dataset from their work.

282 References

283 Michelle Bardis, Roozbeh Houshyar, Chanon Chantaduly, Alexander Ushinsky, Justin Glavis-Bloom,
284 Madeleine Shaver, Daniel Chow, Edward Uchio, and Peter Chang (2020). “Deep Learning with
285 Limited Data: Organ Segmentation Performance by U-Net.” *Electronics* 9.8. ISSN: 2079-9292. DOI:
286 10.3390/electronics9081199. URL: <https://www.mdpi.com/2079-9292/9/8/1199>.

- 287 Yunzhi Chen, Wei Chen, Omid Rahmati, Fatemeh Falah, Dominik Kulakowski, Saro Lee, Fatemeh
288 Rezaie, Mahdi Panahi, Aref Bahmani, Hamid Darabi, Ali Torabi Haghghi, and Huiyuan Bian
289 and (2022). “Toward the development of deep learning analyses for snow avalanche releases in
290 mountain regions.” *Geocarto International* 37.25, pp. 7855–7880. DOI: 10.1080/10106049.
291 2021. 1986578. eprint: <https://doi.org/10.1080/10106049.2021.1986578>. URL:
292 <https://doi.org/10.1080/10106049.2021.1986578>.
- 293 Bahram Choubin, Moslem Borji, Amir Mosavi, Farzaneh Sajedi-Hosseini, Vijay P. Singh, and
294 Shahaboddin Shamshirband (2019). “Snow avalanche hazard prediction using machine learning
295 methods.” *Journal of Hydrology* 577, p. 123929. ISSN: 0022-1694. DOI: <https://doi.org/10.1016/j.jhydrol.2019.123929>. URL: <https://www.sciencedirect.com/science/article/pii/S0022169419306493>.
- 296 National Research Council, Division on Engineering, Physical Sciences, Commission on Engineering,
297 Technical Systems, and Committee on Ground Failure Hazards Mitigation Research (1990). *Snow
300 avalanche hazards and mitigation in the United States*. National Academies Press.
- 301 Federal Office of Topography swisstopo (2022). *SwissALTI3D: High-precision digital elevation model
302 of Switzerland*. <https://www.swisstopo.admin.ch/en/height-model-swissalit3d>.
303 Accessed: 2025-04-25.
- 304 Reinhard Fromm and Christine Schönberger (2022). “Estimating the danger of snow avalanches
305 with a machine learning approach using a comprehensive snow cover model.” *Machine Learning
306 with Applications* 10, p. 100405. ISSN: 2666-8270. DOI: <https://doi.org/10.1016/j.mlwa.2022.100405>. URL: <https://www.sciencedirect.com/science/article/pii/S2666827022000809>.
- 309 E. D. Hafner, P. Barton, R. C. Daudt, J. D. Wegner, K. Schindler, and Y. Bühler (2022). “Automated
310 avalanche mapping from SPOT 6/7 satellite imagery with deep learning: results, evaluation,
311 potential and limitations.” *The Cryosphere* 16.9, pp. 3517–3530. DOI: 10.5194/tc-16-3517-
312 2022. URL: <https://tc.copernicus.org/articles/16/3517/2022/>.
- 313 Huimin Huang, Lanfen Lin, Ruofeng Tong, Hongjie Hu, Qiaowei Zhang, Yutaro Iwamoto, Xianhua
314 Han, Yen-Wei Chen, and Jian Wu (2020). “UNet 3+: A Full-Scale Connected UNet for Medical
315 Image Segmentation.” *ICASSP 2020 - 2020 IEEE International Conference on Acoustics, Speech
316 and Signal Processing (ICASSP)*, pp. 1055–1059. DOI: 10.1109/ICASSP40776.2020.9053405.
- 317 A. Pozdnoukhov, R.S. Purves, and M. Kanevski (2008). “Applying machine learning meth-
318 ods to avalanche forecasting.” *Annals of Glaciology* 49, pp. 107–113. DOI: 10.3189/172756408787814870.
- 320 Jürg Schweizer, J. Bruce Jamieson, and Martin Schneebeli (2003). “Snow avalanche formation.”
321 *Reviews of Geophysics* 41.4. DOI: <https://doi.org/10.1029/2002RG000123>. eprint:
322 <https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2002RG000123>. URL:
323 <https://agupubs.onlinelibrary.wiley.com/doi/abs/10.1029/2002RG000123>.
- 324 Peyman Yariyan, Ebrahim Omidvar, Foad Minaei, Rahim Ali Abbaspour, and John P. Tiefenbacher
325 (Mar. 2022). “An optimization on machine learning algorithms for mapping snow avalanche
326 susceptibility.” *Natural Hazards* 111.1, pp. 79–114. ISSN: 1573-0840. DOI: 10.1007/s11069-
327 021-05045-5. URL: <https://doi.org/10.1007/s11069-021-05045-5>.